

MGMT 467 | Final Project Report | Team 8 | Artemii Chirkov, Kanan Gurbanov, Brandon Moss

Architecture & Ops

This report documents the overall architecture of the project and how to spin up / tear down the data and modeling pipeline.

High-Level Architecture

At a high level, the system combines **batch data** from Kaggle with **streaming data** from a public earthquake API (<https://www.kaggle.com/datasets/usgs/earthquake-database>), then trains a BQML model and surfaces KPIs in Looker Studio.

1. Batch Path – Kaggle → GCS → BigQuery

- **Source**
 - Kaggle earthquakes CSV file from this link:
<https://www.kaggle.com/datasets/usgs/earthquake-database>
- **Storage & raw table**
 - Raw file stored in GCS (e.g., gs://mgmt-467-471119-opensky/earthquakes_raw/).
 - Loaded into a raw BigQuery table mgmt-467-471119.database.earthquakes_raw.
- **Curation**
 - A SQL transformation builds mgmt-467-471119.database.earthquakes_batch which:
 - Casts magnitude, depth, latitude, and longitude to numeric types.
 - Normalizes timestamps into a single TIMESTAMP column.
 - Enforces simple data-quality rules (e.g. $0 \leq \text{depth_km} \leq 700$, magnitude not NULL).
 - Partitions the table by date to support efficient queries.

2. Streaming Path – API → Cloud Function → Pub/Sub → Dataflow → BigQuery

- **Source**
 - Public earthquake API (e.g., USGS) returning JSON for recent events.
- **Cloud Function (2nd gen)**
 - Triggered on a schedule by Cloud Scheduler (e.g., every few minutes).
 - Calls the earthquake API and normalizes JSON into a payload with fields such as event_id, event_time, mag, depth_km, lat, lon, place, and source.
 - Publishes each event to a Pub/Sub topic, for example: earthquake-events.
- **Dataflow streaming job**
 - Consumes messages from the Pub/Sub topic.
 - Applies basic parsing / validation.
 - Writes to a BigQuery streaming table mgmt-467-471119.database.earthquakes_stream.
- **Parsed view for analytics**
 - A SQL view mgmt-467-471119.database.earthquakes_stream_parsed uses JSON_VALUE / SAFE_CAST to turn the JSON payload into typed columns.
 - This view feeds both:
 - The unified earthquakes_live table, and
 - The Looker Studio streaming dashboard.

3. Unified Live Table & BQML Model


- **Unified table**
 - mgmt-467-471119.database.earthquakes_live combines:
 - Historical curated batch data from earthquakes_batch.
 - Recent streaming rows from earthquakes_stream_parsed.
 - Ensures the model sees a consistent schema across batch + streaming.
- **BQML model**
 - Logistic regression classifier mgmt-467-471119.database.severity_classifier.

- Features: mag, depth_km, lat, lon.
- Label: label_severe = 1 if $\text{mag} \geq 5.5$, else 0.
- Evaluation done via ML.EVALUATE, predictions via ML.PREDICT.

4. Analytics & Dashboards


- **Looker Studio dashboard**

- Data source: earthquakes_stream_parsed.
- KPIs (Last 24 Hours):
 - Earthquakes – COUNT(event_id),
 - Avg Magnitude – AVG(mag),
 - Max Magnitude – MAX(mag).
- Time series:
 - “Earthquakes per Day (Streaming)” – count of events over time.



- **Notebook analytics**

- Additional Plotly visualizations (e.g. magnitude vs depth by severity) for model interpretability and KPI design.



Reproducibility & Operations

This checklist explains how to **spin up** and **tear down** the pipeline in GCP.

Spin-Up Steps

1. Enable core GCP APIs

- BigQuery API
- Cloud Storage API
- Cloud Functions API
- Cloud Scheduler API
- Pub/Sub API
- Dataflow API

2. Create storage and BigQuery datasets

- GCS bucket for raw data, e.g. gs://mgmt-467-471119-opensky.
- BigQuery dataset database in project mgmt-467-471119.

3. Batch ingestion

- Upload the Kaggle earthquakes CSV into the GCS bucket.
- Load the file into database.earthquakes_raw.

- Run the SQL script that builds the curated, partitioned database.earthquakes_batch with:
 - Correct types,
 - Timestamp normalization,
 - Data-quality filters (e.g., depth and magnitude rules).

4. Streaming ingestion

- Create Pub/Sub topic earthquake-events.
- Deploy the 2nd-gen Cloud Function that:
 - Calls the earthquake API,
 - Publishes normalized JSON to the earthquake-events topic.
- Configure a Cloud Scheduler job to trigger the function on a fixed schedule.
- Start a Dataflow streaming job using the Pub/Sub → BigQuery template:
 - Input: earthquake-events topic,
 - Output: database.earthquakes_stream.
- Verify that rows appear in database.earthquakes_stream_parsed with recent timestamps.

5. Model training & scoring

- Run the CREATE OR REPLACE MODEL statement for database.severity_classifier.
- Use ML.EVALUATE to check metrics.
- Optionally schedule a query that periodically scores new events in earthquakes_live and writes predictions / probabilities to a table for further reporting.

6. Dashboards & notebooks

- In Looker Studio, connect to mgmt-467-471119.database.earthquakes_stream_parsed as a data source and recreate the 3 KPIs + time series.

- Open this notebook and run all cells to regenerate tables, model evaluation, and Plotly visualizations.

Tear-Down Steps

1. Pause streaming

- Stop or drain the Dataflow streaming job.
- Disable or delete the Cloud Scheduler job.
- Disable or delete the Cloud Function if you no longer want to call the API.

2. Clean up GCP resources

- Optionally delete the Pub/Subtopic earthquake-events.
- Optionally remove any temporary GCS staging buckets used by Dataflow / Cloud Functions.

3. BigQuery cleanup (optional)

- If cost is a concern, drop or expire:
 - database.earthquakes_stream,
 - database.earthquakes_batch,
 - database.earthquakes_live,
 - Any temporary scoring tables.

Documenting these steps makes the project **reproducible** for another person (or a future version of ourselves) and shows how the end-to-end pipeline can be deployed and safely shut down.